

Московский Государственный технический университет  
имени Н. Э. Баумана



Лабораторная работа №2 по курсу:  
«Технология машинного обучения»

Работу выполнил студент группы ИУ5-63  
Федорова Антонина\_\_\_\_\_

Работу проверил:  
Гапанюк Ю.Е.\_\_\_\_\_

Москва  
2019

## **Задание:**

### **Часть 1.**

Выполните первое демонстрационное задание "demo assignment" под названием "Exploratory data analysis with Pandas" со страницы курса

<https://mlcourse.ai/assignments>

Условие задания - [https://nbviewer.jupyter.org/github/Yorko/mlcourse\\_open/blob/master/jupyter\\_english/assignments\\_demo/assignment01\\_pandas\\_uci\\_adult.ipynb?flush\\_cache=true](https://nbviewer.jupyter.org/github/Yorko/mlcourse_open/blob/master/jupyter_english/assignments_demo/assignment01_pandas_uci_adult.ipynb?flush_cache=true)

Набор данных можно скачать здесь - <https://archive.ics.uci.edu/ml/datasets/Adult>

Пример решения задания - <https://www.kaggle.com/kashnitsky/a1-demo-pandas-and-uci-adult-dataset-solution>

### **Часть 2.**

Выполните следующие запросы с использованием двух различных библиотек - **Pandas** и **PandaSQL**:

- один произвольный запрос на соединение двух наборов данных
- один произвольный запрос на группировку набора данных с использованием функций агрегирования

**Текст программы с примерами выполнения программы:**



```
! pip3 install pandasql
```

```
↳ Collecting pandasql
```

```
  Downloading https://files.pythonhosted.org/packages/6b/c4/ee4096ffa2eeec
```

```
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-pack
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.6/dist-pac
```

```
Requirement already satisfied: sqlalchemy in /usr/local/lib/python3.6/dist
```

```
Requirement already satisfied: python-dateutil>=2 in /usr/local/lib/python
```

```
Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dis
```

```
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/dist-p
```

```
Building wheels for collected packages: pandasql
```

```
  Building wheel for pandasql (setup.py) ... done
```

```
  Stored in directory: /root/.cache/pip/wheels/53/6c/18/b87a2e5fa8a82e9c02
```

```
Successfully built pandasql
```

```
Installing collected packages: pandasql
```

```
Successfully installed pandasql-0.7.3
```

## ▼ Часть 1

In this task you should use Pandas to answer a few questions about the Adult dataset. (You don't have to download the data – it's already in the repository).

Choose the answers in the web-form.

Unique values of all features (for more information, please see the links above):

- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.
- salary: >50K,<=50K

```
import numpy as np
import pandas as pd
pd.set_option('display.max.columns', 100)
# to draw pictures in jupyter notebook
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
# we don't like warnings
# you can comment the following 2 lines if you'd like to
import warnings
warnings.filterwarnings('ignore')
from google.colab import drive
```

```
drive.mount("/content/gdrive", force_remount=True)
```

➞ Go to this URL in a browser: <https://accounts.google.com/o/oauth2/auth?cli>

Enter your authorization code:

.....

Mounted at /content/gdrive

```
#Загружаю данные с гугл диска
```

```
data = pd.read_csv('/content/gdrive/My Drive/adult.data.csv')
```

```
data.head()
```

➞

	age	workclass	fnlwgt	education	education- num	marital- status	occupation	rela
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	N
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	N
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	

1. How many men and women (sex feature) are represented in this dataset?

```
data['sex'].value_counts()
```

```

↳ Male      21790
   Female    10771
   Name: sex, dtype: int64

```

## 2. What is the average age (age feature) of women?

```
data.loc[data['sex'] == 'Female', ['age']].mean()
```

```

↳ age      36.85823
   dtype: float64

```

## 3. What is the percentage of German citizens (native-country feature)?

```

#Всего граждан из Германии
(data['native-country'] == 'Germany').sum()

```

```
↳ 137
```

```
(data['native-country'] == 'Germany').sum()/data.shape[0]
```

```
↳ 0.004207487485028101
```

## 4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (salary feature) and those who earn less than 50K per year?

```

poor_mean = data.loc[data['salary'] == '<=50K', 'age'].mean()
poor_std = data.loc[data['salary'] == '<=50K', 'age'].std()
rich_mean = data.loc[data['salary'] == '>50K', 'age'].mean()
rich_std = data.loc[data['salary'] == '>50K', 'age'].std()
print('Средний возраст бедных = {0} +- {1}'.format(round(poor_mean), round(poor_std)))
print('Средний возраст богатых = {0} +- {1}'.format(round(rich_mean), round(rich_std)))

```

```

↳ Средний возраст бедных = 37 +- 14.0
   Средний возраст богатых = 44 +- 10.5

```

## 6. Is it true that people who earn more than 50K have at least high school education? (education – Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)

```
data.loc[data['salary'] == '>50K', 'education'].value_counts()
```

```

[>] Bachelors      2221
     HS-grad      1675
     Some-college  1387
     Masters      959
     Prof-school   423
     Assoc-voc     361
     Doctorate     306
     Assoc-acdm    265
     10th          62
     11th          60
     7th-8th       40
     12th          33
     9th           27
     5th-6th       16
     1st-4th        6
     Name: education, dtype: int64

```

**7. Display age statistics for each race (race feature) and each gender (sex feature). Use `groupby()` and `describe()`. Find the maximum age of men of Amer-Indian-Eskimo race.**

```

for (race, sex), sub_data in data.groupby(['race', 'sex']):
    print("Race: {0}, sex: {1}".format(race, sex))
    print(sub_data['age'].describe())

```

```

[>] Race: Amer-Indian-Eskimo, sex: Female
     count      119.000000
     mean       37.117647
     std        13.114991
     min        17.000000
     25%        27.000000
     50%        36.000000
     75%        46.000000
     max        80.000000
     Name: age, dtype: float64
     Race: Amer-Indian-Eskimo, sex: Male
     count      192.000000
     mean       37.208333
     std        12.049563
     min        17.000000
     25%        28.000000
     50%        35.000000
     75%        45.000000
     max        82.000000
     Name: age, dtype: float64
     Race: Asian-Pac-Islander, sex: Female
     count      346.000000
     mean       35.089595
     std        12.300845
     min        17.000000
     25%        25.000000

```

```

50%      33.000000
75%      43.750000
max       75.000000
Name: age, dtype: float64
Race: Asian-Pac-Islander, sex: Male
count     693.000000
mean       39.073593
std        12.883944
min        18.000000
25%        29.000000
50%        37.000000
75%        46.000000
max        90.000000
Name: age, dtype: float64
Race: Black, sex: Female
count     1555.000000
mean       37.854019
std        12.637197
min        17.000000
25%        28.000000
50%        37.000000
75%        46.000000
max        90.000000
Name: age, dtype: float64
Race: Black, sex: Male
count     1569.000000
mean       37.682600
std        12.882612
min        17.000000
25%        27.000000
50%        36.000000
75%        46.000000
-----

```

**8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (marital-status feature)? Consider as married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.**

```

data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].isin(['Never-married',
                                       'Separated',
                                       'Divorced',
                                       'Widowed'])), 'salary'].value_counts()

```

```

↳  <=50K      7552
   >50K        697
   Name: salary, dtype: int64

```



```
data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].str.startswith('Married')), 'salary'].value_counts()

[>]  <=50K      7576
      >50K      5965
      Name: salary, dtype: int64
```

**9. What is the maximum number of hours a person works per week (hours-per-week feature)? How many people work such a number of hours, and what is the percentage of those who earn a lot (>50K) among them?**

```
max_time = data['hours-per-week'].max()
print('Максимальное число часов работы: {0}'.format(max_time))
print('Число людей, работающих {0} часов в неделю и их зарплаты:'.format(max_time))
print(data.loc[data['hours-per-week'] == max_time, 'salary'].value_counts())
perc = ((data['hours-per-week'] == max_time) & (data['salary'] == '>50K')).sum()
print('Процент людей, которые работают по {0} часов и получают зарплату более 50K: {0}'.format(perc))
```

```
[>]  Максимальное число часов работы: 99
      Число людей, работающих 99 часов в неделю и их зарплаты:
      <=50K      60
      >50K      25
      Name: salary, dtype: int64
      Процент людей, которые работают по 99 часов и получают зарплату более 50K: 0.020202
```

**10. Count the average time of work (hours-per-week) for those who earn a little and a lot (salary) for each country (native-country). What will these be for Japan?**

```
for (salary, native_country), sub_data in data.groupby(['salary', 'native-country']):
    print("native-country: {0}, salary: {1}".format(native_country, salary))
    print(round(sub_data['hours-per-week'].mean(), 2))
```

```
[>]  native-country: ?, salary: <=50K
      40.16
      native-country: Cambodia, salary: <=50K
      41.42
      native-country: Canada, salary: <=50K
      37.91
      native-country: China, salary: <=50K
      37.38
      native-country: Columbia, salary: <=50K
      38.68
      native-country: Cuba, salary: <=50K
      37.99
      native-country: Dominican-Republic, salary: <=50K
      42.34
      native-country: Ecuador, salary: <=50K
      38.04
      native-country: El-Salvador, salary: <=50K
      36.03
```

```

native-country: England, salary: <=50K
40.48
native-country: France, salary: <=50K
41.06
native-country: Germany, salary: <=50K
39.14
native-country: Greece, salary: <=50K
41.81
native-country: Guatemala, salary: <=50K
39.36
native-country: Haiti, salary: <=50K
36.33
native-country: Holand-Netherlands, salary: <=50K
40.0
native-country: Honduras, salary: <=50K
34.33
native-country: Hong, salary: <=50K
39.14
native-country: Hungary, salary: <=50K
31.3
native-country: India, salary: <=50K
38.23
native-country: Iran, salary: <=50K
41.44
native-country: Ireland, salary: <=50K
40.95
native-country: Italy, salary: <=50K
39.62
native-country: Jamaica, salary: <=50K
38.24
native-country: Japan, salary: <=50K
41.0
native-country: Laos, salary: <=50K
40.38
native-country: Mexico, salary: <=50K
40.0
native-country: Nicaragua, salary: <=50K
36.09
native-country: Outlying-US(Guam-USVI-etc), salary: <=50K
41.86
native-country: Peru, salary: <=50K

```

## ▼ Часть 2

```

user_usage = pd.read_csv('/content/gdrive/My Drive/user_usage.csv')
android_devices = pd.read_csv('/content/gdrive/My Drive/android_devices.csv')
user_device = pd.read_csv('/content/gdrive/My Drive/user_device.csv')

```

```
user_usage.head()
```

	<b>outgoing_mins_per_month</b>	<b>outgoing_sms_per_month</b>	<b>monthly_mb</b>	<b>use_id</b>
0	21.97	4.82	1557.33	22787
1	1710.08	136.88	7267.55	22788
2	1710.08	136.88	7267.55	22789
3	94.46	35.17	519.12	22790
4	71.59	79.26	1557.33	22792

```
android_devices.head()
```

	<b>Retail</b>	<b>Branding</b>	<b>Marketing Name</b>	<b>Device</b>	<b>Model</b>
0		NaN	NaN	AD681H	Smartfren Andromax AD681H
1		NaN	NaN	FJL21	FJL21
2		NaN	NaN	T31	Panasonic T31
3		NaN	NaN	hws7721g	MediaPad 7 Youth 2
4		3Q	OC1020A	OC1020A	OC1020A

```
user_device.head()
```

	<b>use_id</b>	<b>user_id</b>	<b>platform</b>	<b>platform_version</b>	<b>device</b>	<b>use_type_id</b>
0	22782	26980	ios	10.2	iPhone7,2	2
1	22783	29628	android	6.0	Nexus 5	3
2	22784	28473	android	5.1	SM-G903F	1
3	22785	15200	ios	10.2	iPhone7,2	3
4	22786	28239	android	6.0	ONE E1003	1

## ▼ Pandas

```
full1 = pd.merge(user_usage,
                  user_device,
                  on='use_id',
                  how = 'left')
```

```
full1.head()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	use
0	21.97	4.82	1557.33	22787	129
1	1710.08	136.88	7267.55	22788	28
2	1710.08	136.88	7267.55	22789	28
3	94.46	35.17	519.12	22790	29
4	71.59	79.26	1557.33	22792	28

```
full1.tail()
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	u
235	260.66	68.44	896.96	25008	
236	97.12	36.50	2815.00	25040	
237	355.93	12.37	6828.09	25046	
238	632.06	120.46	1453.16	25058	
239	488.70	906.92	3089.85	25220	

Можно заметить, что появляются нулевые поля так как я использовала left join. То есть в итоговую таблицу помещаются те значения, которые есть в первой таблице, но их нет во второй таблице.

```
full1.agg(
    {'outgoing_mins_per_month' : ['sum', 'min', 'mean', 'max'],
     'outgoing_sms_per_month' : ['min', 'max', 'sum', 'mean'],
     'platform_version' : ['min', 'max', 'sum', 'mean']}
)
```

	outgoing_mins_per_month	outgoing_sms_per_month	platform_version
max	1816.630000	906.920000	10.100000
mean	274.559167	98.968292	5.554717
min	0.500000	0.250000	4.100000
sum	65894.200000	23752.390000	883.200000

Вывела для трех значений максимум, среднее, минимум, сумму (с помощью функции агрегации)

## ▼ PandaSQL

```
import pandasql as ps

def example_query(full1):
    simple_query = '''
        SELECT
            outgoing_mins_per_month,
            monthly_mb,
            platform
        FROM full1
        where 1=1
        and monthly_mb >= 100
        --and platform = 'ios'
        LIMIT 15
    '''
    return ps.sqldf(simple_query, locals())
```

```
example_query(full1)
```

→	outgoing_mins_per_month	monthly_mb	platform
0	21.97	1557.33	android
1	1710.08	7267.55	android
2	1710.08	7267.55	android
3	94.46	519.12	android
4	71.59	1557.33	android
5	71.59	1557.33	android
6	71.59	519.12	android
7	71.59	519.12	android
8	30.92	3114.67	android
9	69.80	25955.55	android
10	554.41	3114.67	android
11	189.10	519.12	android
12	283.30	15573.33	android
13	324.34	519.12	android
14	797.06	519.12	android

```
def aggr_query(full1):  
    aggr_query = '''  
        SELECT  
            count(*),  
            platform  
        FROM full1  
        group by platform  
        LIMIT 15  
    '''  
    return ps.sqldf(aggr_query, locals())
```

```
aggr_query(full1)
```

	count(*)	platform
0	81	None
1	157	android
2	2	ios

```
def join_query(user_usage, user_device):
    join_query = '''
    select *
    from user_usage t0
    join user_device t1
    on t0.use_id = t1.use_id'''

    return ps.sqldf(join_query)
```

```
join_query(user_usage, user_device)
```

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	u
0	21.97	4.82	1557.33	22787	
1	1710.08	136.88	7267.55	22788	
2	1710.08	136.88	7267.55	22789	
3	94.46	35.17	519.12	22790	
4	71.59	79.26	1557.33	22792	
5	71.59	79.26	1557.33	22793	
6	71.59	79.26	519.12	22794	
7	71.59	79.26	519.12	22795	
8	30.92	22.77	3114.67	22799	
9	69.80	14.70	25955.55	22801	
10	554.41	150.06	3114.67	22804	
11	189.10	24.08	519.12	22805	
12	283.30	107.47	15573.33	22806	
13	324.34	92.52	519.12	22808	
14	797.06	7.67	519.12	22813	

<b>15</b>	797.06	7.67	15573.33	22814
<b>16</b>	797.06	7.67	15573.33	22815
<b>17</b>	797.06	7.67	15573.33	22816
<b>18</b>	797.06	7.67	15573.33	22817
<b>19</b>	78.80	327.33	10382.21	22819
<b>20</b>	78.80	327.33	15573.33	22820
<b>21</b>	78.80	327.33	15573.33	22822
<b>22</b>	164.10	192.64	3114.67	22823
<b>23</b>	208.26	91.76	5191.12	22824
<b>24</b>	681.44	47.35	1271.39	22829
<b>25</b>	324.27	91.50	519.12	22830